

A. Project Overview

- **Goal:** What question are you answering?

The question I am answering through this project is looking at how the death rate changes throughout the years. I'm looking at the death rate for each year and each country, both genders, and for all ages under age group (since it was divided into multiple age groups).

- **Dataset:** Source, size (link if too large for GitHub).

I got my dataset from [tableau](#) where they have free public datasets for analysis. The specific data I used for my project is: [a public health data set](#).

B. Data Processing

- How you **loaded** it into Rust.
- Any **cleaning** or **transformations** applied.

I first downloaded the dataset from the website (downloading the linked file called "Global Burden of Disease" under Free Health Data Sets). I then added this dataset into my "proj" directory inside of my ds210project folder so that my Rust functions would be able to access it. I renamed the csv to a shorter and simpler name called "health.csv". I didn't clean the data prior to loading it, I cleaned/filtered the data through my Rust functions.

C. Code Structure

1. Modules

- Purpose of each and rationale for organization.

My Rust project is organized into 4 modules: main.rs, dataframe.rs, linear.rs, and plot.rs which handles different things for my project.

- Dataframe.rs = creates a Dataframe structure from a CSV, handling the reading of the CSV, loading, storing, and also operations like printing and filtering the data.
- Linear.rs = handles training a linear regression model given data taken from the Dataframe given the column names through the linfa and linfa_linear crates. After locating the position of those columns in the dataset, it extracts it and converts it to ndarrays to fit the model.
- Plot.rs = generates a scatter plot given the two selected columns and also puts a fitted regression line through it using the plotters crate to visualize it. The plot gets saved as an image file and the linear regression equation based on this model also gets printed.

- Main.rs = it runs the whole thing: loading the data from the csv I uploaded, filtering it (age group and gender), training a linear regression model based on the data, and generating a plot of it. It also has tests to verify the code is functioning the way it should.

The project is organized into these different modules so that each can handle a different task to make it more manageable, to fix, and to read.

2. Key Functions & Types (Structs, Enums, Traits, etc)

- For each non-trivial item, restate:
 - **Purpose**
 - **Inputs/Outputs**
 - **Core logic and key components**

For my dataframe.rs:

I have a Dataframe struct which will store the labels, data, and the column types. The ColumnVal enum supports the different types of data in the columns, and for this data it represents one of 3 types: string, integer, or a float.

My read_csv function will load the csv into a Dataframe with the inputs of a file path and a list of expected column types. It reads the labels and rows, parses the values based on type, and then stores it as ColumnVal values. It returns OK(()) if it's successful or an error if not.

My filtering function will return a new Dataframe with just the rows that match the conditions. It takes the column label and a filtering closure as the input and outputs a filtered Dataframe. It gets the index of the specified column, applies the closure to that column, and returns the matching rows.

For my Linear.rs:

My fit_model function will fit a linear regression model to the two columns I specified. It takes the data, labels, x column name, and y column name as inputs. It will output a trained FittedLinearRegression<f64> model by taking the values from the selected x and y column, creating ndarray dataset, and fitting it through linfa linear regression.

For my plot.rs:

My plot_data function will create a scatter plot with a regression line, and print the regression line equation. It takes the data rows, labels, x axis, and y axis column names as inputs. The output is a generated png and returns ok(()) on success. It finds the right columns and filters the data, and uses plotters to do the scatter plot and fit the line. It then calls fit_model to get the slope and intercept for the equation.

3. Main Workflow

- At a glance, how modules/functions interact to produce your results.

For my main.rs:

It loads the dataset through calling `read_csv` function and applies the filter function twice by age and gender (to get the desired values for those columns). It calls `plot_data` to visualize the data & regression line, and also print the equation for it. Main uses the `dataframe.rs` module to process the data, uses the `linear.rs` module to analyze it, and the `plot.rs` to visualize it.

D. Tests

- **cargo test output** (paste logs or provide screenshots).
- For each test: what it checks and why it matters.

I have two tests for this project: one that checks if it reads the csv correctly and one that checks whether the linear regression outputs the expected slope and y-intercept.

For the csv test, it should return `ok(())` meaning there is no error and the labels in the Dataframe match the one in the csv. It matters because it sets a foundation for the other functions. If it doesn't read or label the csv correctly, it could affect the filtering, plotting, and modeling linear regression.

For the linear regression test, it tests if the linear regression model accurately fits the data. To test this one, I made a sample csv called "test_data.csv" where the x and y values are computed based on the equation $y = 5x + 50$. I call the `fit_model` function to train the columns with those x and y values to check if it learns that the slope is approximately 5.0 and y intercept is approximately 50.0. (It is approximate because the floats can have small rounding errors).

My cargo test output:

```
[/opt/app-root/src/ds210project/proj]
● $ cargo test
    Finished `test` profile [unoptimized + debuginfo] target(s) in 0.08s
    Running unittests src/main.rs (target/debug/deps/proj-915b44425a96fd93)

running 2 tests
test tests::it_reads ... ok
test tests::linear_regression_works ... ok

test result: ok. 2 passed; 0 failed; 0 ignored; 0 measured; 0 filtered out; finished in 0.00s
```

E. Results

- All program outputs (screenshots or pasted).
- Interpretation in project context (no need for "groundbreaking" results).

My output for the filtered dataset & plotting the results:


```

$ cargo run
Compiling proj v0.1.0 (/opt/app-root/src/ds210project/proj)
Finished dev profile [unoptimized + debuginfo] target(s) in 1.76s
Running target/debug/proj
Country Code Country Name Year Age Group Sex Number of Deaths Death Rate Per 100,000
AFG Afghanistan 1970 All ages Both 291837 2432.3
AFG Afghanistan 1980 All ages Both 292045 2026.8
AFG Afghanistan 1990 All ages Both 200636 1486.2
AFG Afghanistan 2000 All ages Both 335117 1445.4
AFG Afghanistan 2010 All ages Both 302163 946.3
AGO Angola 1970 All ages Both 138582 2321.3
AGO Angola 1980 All ages Both 158589 2071.1
AGO Angola 1990 All ages Both 191100 1830.7
AGO Angola 2000 All ages Both 201736 1433.1
AGO Angola 2010 All ages Both 161489 835.7
ALB Albania 1970 All ages Both 13716 640.1
ALB Albania 1980 All ages Both 15171 567.8
ALB Albania 1990 All ages Both 16500 502.9
ALB Albania 2000 All ages Both 19308 627
ALB Albania 2010 All ages Both 22659 708.1
AND Andorra 1970 All ages Both 134 679.1
AND Andorra 1980 All ages Both 242 697.2
AND Andorra 1990 All ages Both 415 792.8
AND Andorra 2000 All ages Both 513 796.2
AND Andorra 2010 All ages Both 752 899.9

WSM Samoa 1980 All ages Both 794 503.9
WSM Samoa 1990 All ages Both 983 599.3
WSM Samoa 2000 All ages Both 1046 589.2
YEM Yemen 1970 All ages Both 134529 2205.3
YEM Yemen 1980 All ages Both 127724 1604.4
YEM Yemen 1990 All ages Both 122285 1020.7
YEM Yemen 2000 All ages Both 143644 809.9
YEM Yemen 2010 All ages Both 146303 607.4
ZAF South Africa 1970 All ages Both 241258 1069.2
ZAF South Africa 1980 All ages Both 250676 861.5
ZAF South Africa 1990 All ages Both 253495 688.8
ZAF South Africa 2000 All ages Both 459550 1024.1
ZAF South Africa 2010 All ages Both 516120 1027.4
ZMB Zambia 1970 All ages Both 56326 1359.3
ZMB Zambia 1980 All ages Both 70443 1220.3
ZMB Zambia 1990 All ages Both 115278 1468.6
ZMB Zambia 2000 All ages Both 190937 1868.6
ZMB Zambia 2010 All ages Both 148445 1124.4
ZWE Zimbabwe 1970 All ages Both 61795 1181.7
ZWE Zimbabwe 1980 All ages Both 73253 1003.6
ZWE Zimbabwe 1990 All ages Both 78233 745.7
ZWE Zimbabwe 2000 All ages Both 167896 1336.6
ZWE Zimbabwe 2010 All ages Both 159430 1264.9

y = -9.11x + 19087.49
[/opt/app-root/src/ds210project/proj]

```



- The filtered dataset only has data for all ages from the age group column and both gender from 1970 - 2010 for each of the countries. (there is more printed but in the screenshot I have the beginning and last few)
- The plot_data produces a visualization of the results and also the regression line equation. The slope is negative meaning that the death rate has decreased over the

years. The line in the plot is fitted quite well (going through where the average dots seems to be at) and also shows this decrease in death rates with the downward trend.

F. Usage Instructions

- How to build and run your code.
- Description of any command-line arguments or user interaction in the terminal.
- Include expected runtime especially if your project takes a long time to run.

You will just need to do cargo run to run this program. It runs quickly, probably taking a few seconds.

G. AI-Assistance Disclosure and Other Citations

- Cite any **substantive** ChatGPT/GenAI you used (e.g. screenshot or description).
 - You can skip this for common knowledge/debugging use cases
- For each cited snippet, include your own explanation to show understanding.
- You can also provide links to other sources you found useful that are not "common knowledge"

For plotting:

<https://docs.rs/plotters/latest/plotters/#plotting-in-rust>
<https://chatgpt.com/share/680ee033-e198-8001-b30a-569b0e6ecb0b>

I've used chatgpt to understand how the code to plot a scatter plot works (I took the code from the first link to plot and asked chatgpt to explain it) to which I then adjusted the code to fit my data. I also used it to fix the errors I was having when running the plot_data function. I understood now the different components of plotting the data and how to generate an image file such as .build_cartesian_2d() to define the x and y ranges for the chart, .draw_series() to determine the type you want (there are LineSeries, Circles, etc), .caption() to set a title, and others. To extract my x and y values, I could use match statements to match the enum so I can get the values. I knew that I needed to locate the index of my columns, but wasn't sure if there was a method to do so. Now, I know that I could iterate through my string of labels and through .position() find the position of that column label. When I was running into an error with the way I was passing my points that needed to be plotted into the Circle::new() method, I also used chatgpt for help. I was getting that error because points was a Vec<(f64,f64)> but Circle::new() expects an iterator of points and to turn a vec into an iterator, I could use .into_iter() and then map each of the x and y tuple to a Circle element.

For test functions:

<https://chatgpt.com/share/6812e74d-c408-8001-8fc9-db8e595d869a>

I also used chatgpt for understanding the parameters for my `assert!()` macro to check if my condition is true. I wasn't sure what would go inside the `assert!()` for a function that returns `ok()`. I found out that for those I could use the `is_ok()` method.

<https://chatgpt.com/share/681449d2-26dc-8001-8d37-9195e4a93310>

I ran into an error with floating-point precision and not being able to represent the numbers exactly. In my code, I was using `assert_eq!()` to do the comparison, but it would not work even if it was very close, it had to be exact for it to run successfully. Instead of checking if it's exact, I could check if it's approximately equal to 50.0. I would retrieve the y intercept and slope from the `fit_model` function and get the difference from the expected value, use `.abs()` to get the absolute value of that difference, and see if that value is very small.

I've used code from lecture 35 for plotting the points and the linear regression line. I also used similar code to convert my csv into a dataframe from homework 8. I have also received help at office hours from TAs, CAs, and professors.